

Because there were such a small number of possible variations, I tested my program by putting the actual numbers in the files and running the program. I then looked to see whether the output file matched my expectations for what the program was supposed to do. I tested the following configurations:

- One file had even numbers, one file had odd numbers
- One file had number 1-5, the other file had numbers 6-10
- One file had many more numbers than the other
- The files had duplicates of one another (eg, file 1 has 1123 and file 2 has 1156)

In the future, it would be nice to automate testing like this to make sure the algorithm I developed works correctly.